



Parallel Design Patterns Coursework2 Report

B156924

April 10, 2020

Contents

1	Introduction	1
2	Program design	1
2.1	Master/Worker strategy	1
2.2	Actor pattern	2
2.2.1	Master actor	2
2.2.2	Control actor	4
2.2.3	Cell actor	4
2.2.4	Squirrel actor	6
3	Results correctness	7
3.1	Method 1	9
3.2	Method 2	10
4	Conclusions	11

1 Introduction

Simulation calculations can help scientists take advantage of the powerful computing capabilities of supercomputers to analyze and predict complex problems in the real world. This technology has applied to many scientific field, such as infectology, biology, physics and astronomy. In the past few months, the COVID-19 has been spreading among many countries. A lot of data scientists have simulated the spread of the virus and predict the trends of infection and death numbers, which provides useful information to government's actions. At the same time, many scholars have demonstrated to the public the positive effects of home isolation on mitigating the epidemic through scientific simulation.

Different simulations are applicable to different program structures and patterns. For the case in this coursework, it is to help biologists to simulate the spread of disease through the squirrel population for discovering information about protecting the animals. As mentioned in the previous report, the Actor algorithm pattern is the most suitable one among all patterns since it can easily simulate the behavior of individual squirrels in the natural environment and the interaction between squirrels and the environment.

Combined with Master/Worker implementation strategy, the model can be easily constructed. MPI(Message-passing Interface) is applied for message passing between actors. This model is written in C and can be compiled with GCC 4.8.5. The code is tested on 7 computing nodes(each node has 36 cores) on the Cirrus platform with Linux system.

This report is organized as follows. In chapter 2 the design of this simulation program is illustrated. Chapter 3 presents the results of the simulation including correctness test and output analysis. Chapter 4 presents the performance analysis for this program. Finally, in chapter 5, the overall performance of this program is discussed, followed with ideas about further improvement.

2 Program design

In this chapter, the design of this program is described. The Master/Worker strategy is implemented to set a master process to allocate, manage and kill other worker processes. The provided file *pool.c* makes it easy to achieve this strategy. For the Actor pattern, there are 4 actor in this program: *Master*, *Control*, *Cell* and *Squirrel*. The following sections explain the details of their work.

2.1 Master/Worker strategy

Master/Worker strategy is used for tasks that are not tightly coupled. The master process can allocate processes for the coming work to execute and remind them to stop when

the program is going to terminate. The processes can be reused when the previous work is done, so this is conducive to the workload balance.

The file *pool.c* provides a dynamic process pool, and there are a bunch of functions designed for this pool. For this simulation, the master process is used for initializing *processPoolInit()* and shutting down *processPoolFinalise()* the pool, and supervising *masterPool()* each process's work status. For each worker process, it is able to request a new process *startWorkerProcess()* from the master, end its work to sleep *workerSleep()* and request to shut down *shutdownPool()* the pool.

By calling these functions, the simulation program can achieve process allocation and management.

2.2 Actor pattern

In the Actor pattern, actors can perform their own computational work when messages haven't arrived. They are also able to create a new actor, die and send messages to others. These properties are suitable for the biological model.

For this simulation, there are four actors designed: *Master*, *Control*, *Cell* and *Squirrel*. Figure 1 is the interaction diagram of these actors. It is worth noting that the master actor only allocates a process for the control actor, and the control actor allocates the cell actors and squirrel actors including initialized healthy and infected squirrels. The control actor also has a clock function, it will send messages to cell actors when a new month begins and receive cells' updated information. The squirrel actors don't know the clock, and they spend their lives on interacting with the cell actors and the control actor and judging the healthy and reproduce conditions.

In the next sections, the work for each type of actor is explained.

2.2.1 Master actor

In the main function, by calling the *processPoolInit()* function, the master process is created with its status code equaling 2. The work for the master actor is very simple. The first thing is to allocate a process to the control actor, the other thing is to continually check the status of each actor. If some actor calls *shutdownPool()* to request the master to terminate the program, the master will break the check loop and execute *processPoolFinalise()* function, which will send each actors messages to stop their works and to sleep.

If the status code is 1, it means all worker processes. The code in this part is executed in every worker process. There is a while loop designed here keeping receiving messages from other actors who have requested a new process. The message content is the type of actors corresponding to different work: *CONTROL_ACTOR*, *CELL_ACTOR*, *SQUIRREL_ACTOR*, *NEW_SQUIRREL_ACTOR* and *INFECTED_SQUIRREL_ACTOR*.

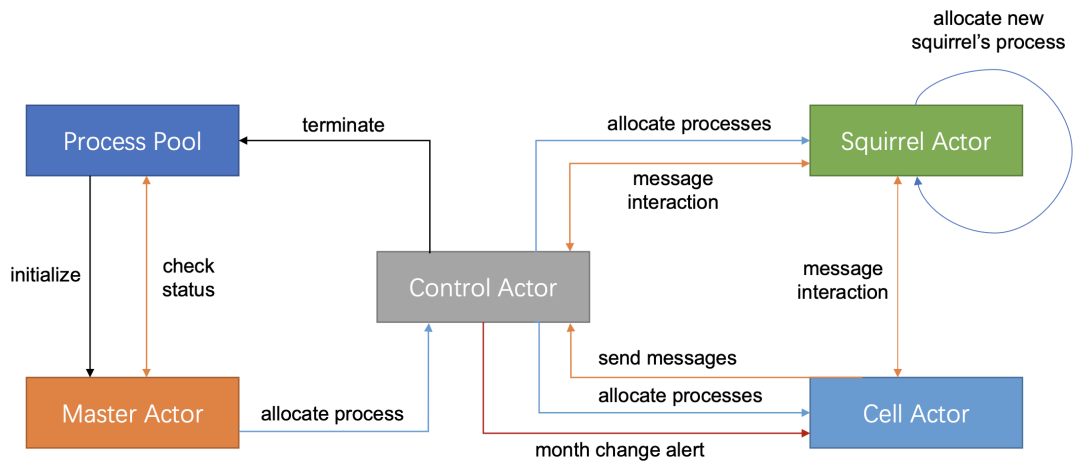


Figure 1: This is the interaction diagram of the actors in the simulation. The process pool is the combination of a bunch of functions that support Master/Worker structure. Therefore, we can regard it as a black box. In the beginning, the master actor is initialized by the process pool. The master actor allocates a new process as the control actor, and then it will continually check the changes of other actors' statuses. As for the control actor, it allocates the processes to all cell actors and squirrel actors. The control actor needs to alert all cell actors per month to update their information and send it to the control actor. The squirrel actor always interacts with cell actors and the control actor and sometimes can born a new squirrel and allocate its process.

When an actor wants to request for a new process, it needs to call *startWorkerProcess()* first and send the type of this actor to its process rank. In this way, a new actor is set up.

2.2.2 Control actor

The control actor is mainly responsible for allocating initialized cell actors and healthy and infected squirrel actors, playing as a clock to interact with cell actors when the month changes and recording the number of alive, infected and dead squirrels for each month. Figure 2 is the workflow of the control actor.

For initializing actors, as mentioned before, the control actor can allocate a process for a new actor and send the work type to this process.

The next step is being a clock and a recorder. when a month passes, the control actor will send messages to each cell actor to remind them of updating their *populationInflux* and *infectionLevel* and then these cells are required to send back their information to the control actor. All cells' information is received with a designed *MPI_Type_vector* and stored in a 2-dimensional array in order.

During each month, the control actor keeps receiving messages from squirrels through *MPI_Iprobe()* detection. There are three types of messages: a new squirrel is born, a healthy squirrel is infected, and an infected squirrel has died. The control actor records the information and displays this information when the month changes.

It is worth noting that when a new month starts, the control actor will judge whether the current squirrel number is 0 or over 200. If all squirrels die, the program needs to be terminated. The control actor will break the loop and call *shutdownPool()*. If the squirrel number is over 200, it means the program may have bugs, therefore, an *assert* function is called to terminate the program straightly.

When the overall simulation time has passed, the control actor will call *shutdownPool()* to tell the master to terminate the program.

2.2.3 Cell actor

The cell actor's work is to record the visits of healthy and infected squirrels, compute *populationInflux* and *infectionLevel* per month and send the updated information to the control actor per month. To store a single month's *populationInflux* and *infectionLevel*, two int arrays *month_populationInflux*, *month_infectionlevel* are created with the size of model months. Figure 3 is the workflow of a cell actor.

When the simulation starts, the cell actor waits for the control actor's first-month message to start working. In this way, the cells won't interact with squirrels immediately after being created, and the accuracy of time can be guaranteed.

Once it receives the month signal, it will jump into a loop. For each iteration, the cell will check whether it needs to terminate. Then a *MPI_Iprobe()* will detect the coming

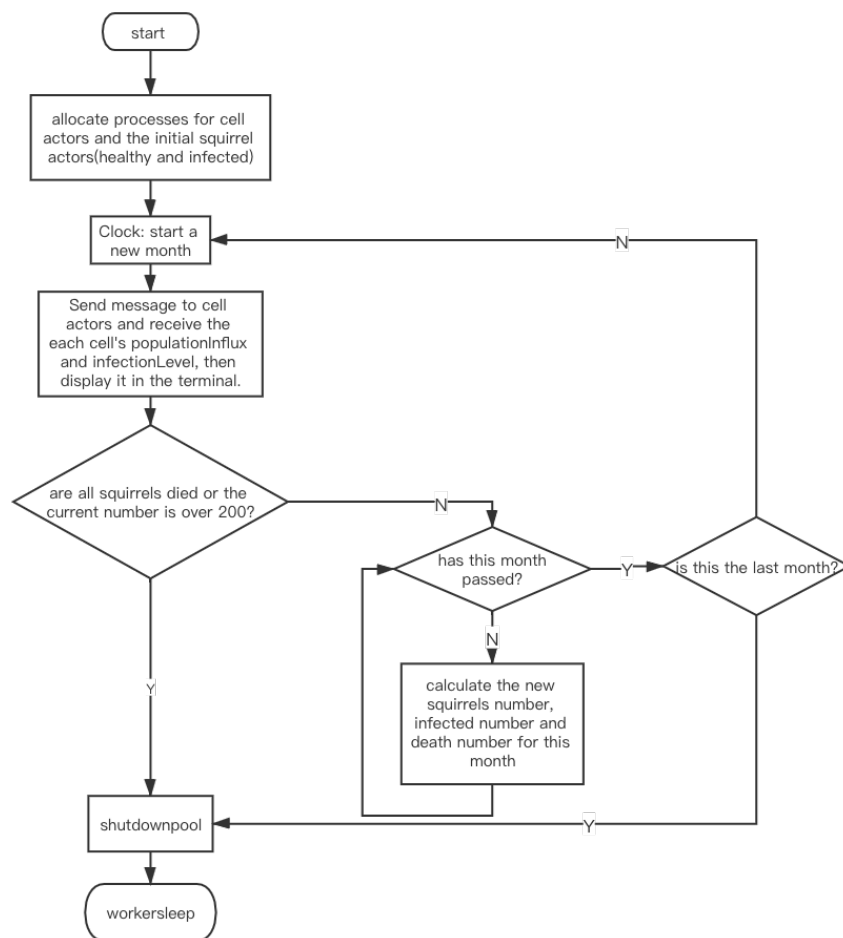


Figure 2: This is the workflow of the control actor.

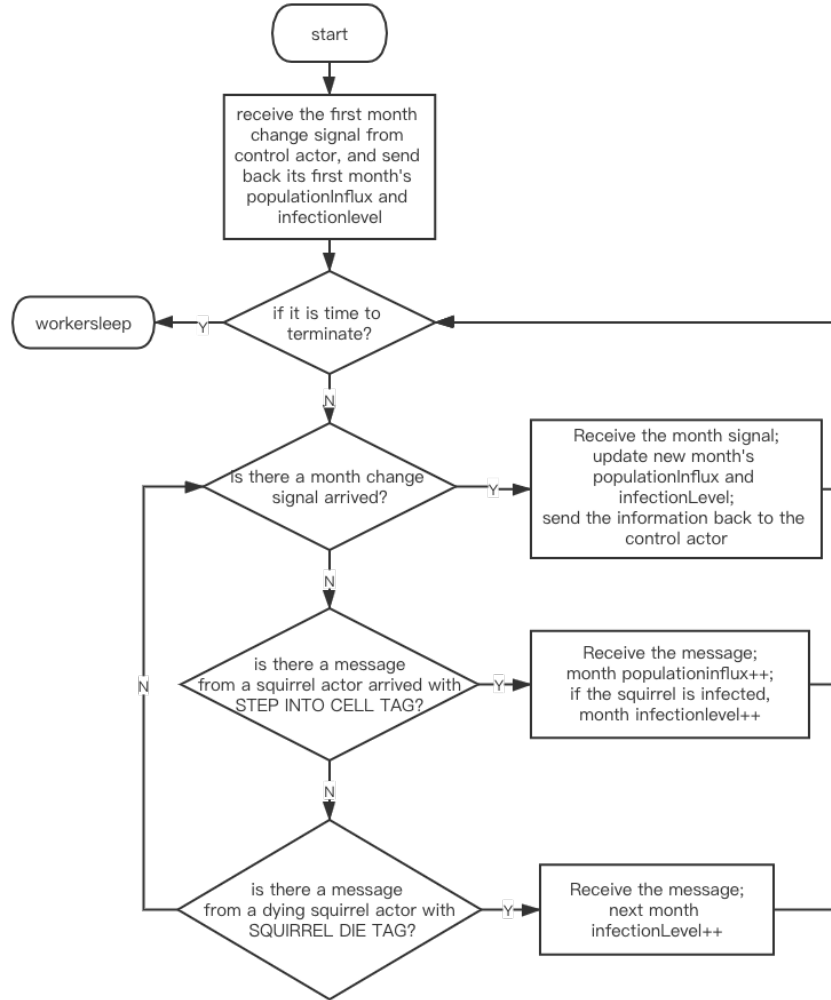


Figure 3: This is the workflow of a cell actor.

messages. If its tag is 'MONTH_ALERT_TAG', the cell needs to update its information and send it back to the control. If it is a 'STEP_INTO_CELL_TAG', it means a squirrel has stepped into this cell. The cell will record this visit to its *month_populationInflux*. If the squirrel is infected, *month_infectionlevel* will add 1. Finally, if its tag is 'SQUIRREL_DIE_TAG', it means that a squirrel is about to die in this cell, and the cell will update its next month's infectionLevel in advance since the virus will remain for 2 months after the squirrel dies.

The cell will keep working in this loop until it is asked to terminate.

2.2.4 Squirrel actor

The squirrel actor simulates the real squirrel's behavior in nature. It will keep jumping between cells and can possibly born a new squirrel, get infected and die. Figure 4 is the

workflow of a squirrel actor.

Since the squirrel sends many messages to the control actor and cell actors, the deadlock can easily appear here when the program is terminated while the squirrel is sending a message. Therefore, before each send, the squirrel will call *workerSleep()* to check whether it is asked to stop working. Besides, the send buffer is allocated with a big size to ensure that the squirrel won't run out of buffer when using the *Bsend()*.

When the squirrel steps into a cell, it will send its information to the cell and receive the cell's *populationInflux* and *infectionLevel*. In order to store the cell information for each step properly, a queue is designed for the squirrel's *populationInflux* trajectory and *infectionLevel* trajectory. The size of the queue is 50. When the queue is full, the earliest step's data will be popped and the newest data will be pushed into the queue. Besides, the mean of all data can be calculated easily and all elements can be printed. Such a design can avoid the wasted storage of useless data and more convenient to calculate.

After storing the cell's data, the squirrel is supposed to test whether it is infected. If the squirrel is infected, the *infected_step* int variable will keep recording its steps since the first step when it is tested infected. If the *infected_step* is more than 50, the squirrel will have 1/6 possibility to death for each next step. When the squirrel is tested dead in a cell, the dying squirrel will send a message with *SQUIRREL_DIE_TAG* to the corresponding cell and a message to the control actor to record its death. After that, it will break the loop and call *workerSleep()* to finish its life.

The final test is whether the squirrel will bear a baby in this step. If it is true, firstly the squirrel needs to ask the control actor whether the current squirrels are more than 200. If not, the squirrel is allowed to allocate a new process for the baby and deliver its position to its baby.

It is worth mentioning that a *squirrel_stall()* function is designed for adjusting the moving speed of squirrels. When the supercomputer computes quickly, all squirrels may die in 2-3 months, the trends of born and death can hardly be analyzed. Therefore, this function can solve this problem.

3 Results correctness

The output of the simulation can be displayed in the terminal. For each month, the alive number, infected number, healthy number, last month death and born number, and the *populationInflux* and *infectionLevel* for each cell are presented with good format. With the same parameters, the results vary in every execution but the general trends are similar. Therefore, the correctness of this simulation needs to be tested to guarantee its accuracy. There are two ways to test its correctness.

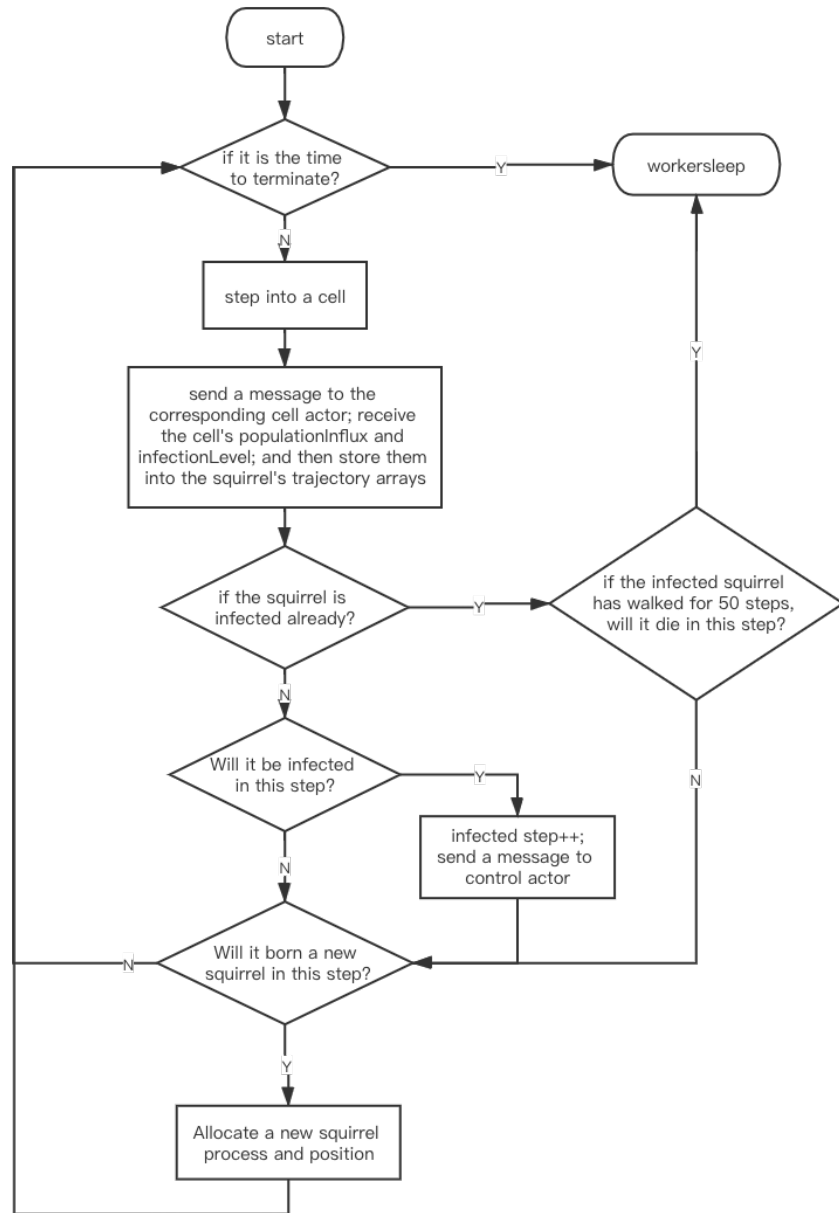


Figure 4: This is the workflow of a squirrel actor.

SIMULATION STARTS.

MONTH: 1	ALIVE: 38	INFECTED: 4	HEALTHY: 34	LAST MONTH DEATH: 0	LAST MONTH BORN: 0											
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
PopulationInflux	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
InfectionLevel	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
MONTH: 2	ALIVE: 38	INFECTED: 4	HEALTHY: 34	LAST MONTH DEATH: 0	LAST MONTH BORN: 0											
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
PopulationInflux	107	102	113	100	100	82	104	104	109	93	95	113	106	97	105	90
InfectionLevel	14	10	14	13	11	11	7	7	11	7	11	14	9	6	7	4
MONTH: 3	ALIVE: 34	INFECTED: 7	HEALTHY: 27	LAST MONTH DEATH: 5	LAST MONTH BORN: 1											
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
PopulationInflux	207	188	211	202	202	183	217	206	245	211	205	225	252	196	211	185
InfectionLevel	22	20	27	30	22	27	24	20	28	16	27	31	27	14	17	11
MONTH: 4	ALIVE: 37	INFECTED: 18	HEALTHY: 19	LAST MONTH DEATH: 3	LAST MONTH BORN: 6											
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
PopulationInflux	299	283	320	300	313	302	323	317	366	316	318	331	348	309	309	280
InfectionLevel	49	49	55	49	62	67	61	60	68	54	58	63	60	62	52	50
MONTH: 5	ALIVE: 23	INFECTED: 13	HEALTHY: 10	LAST MONTH DEATH: 17	LAST MONTH BORN: 3											
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
PopulationInflux	273	254	260	285	310	297	294	301	336	320	299	292	336	287	277	278
InfectionLevel	86	82	75	84	97	91	84	98	95	100	80	84	93	98	86	91
MONTH: 6	ALIVE: 19	INFECTED: 16	HEALTHY: 3	LAST MONTH DEATH: 8	LAST MONTH BORN: 4											
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
PopulationInflux	216	215	221	234	259	246	233	255	266	243	245	239	233	236	216	225
InfectionLevel	78	75	78	88	80	77	81	91	98	89	80	86	83	83	78	84
MONTH: 7	ALIVE: 7	INFECTED: 7	HEALTHY: 0	LAST MONTH DEATH: 13	LAST MONTH BORN: 1											
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
PopulationInflux	153	154	145	167	182	151	159	170	167	173	164	162	162	154	149	153
InfectionLevel	59	63	77	64	62	60	72	62	76	65	74	72	55	66	64	57
MONTH: 8	ALIVE: 3	INFECTED: 3	HEALTHY: 0	LAST MONTH DEATH: 5	LAST MONTH BORN: 1											
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
PopulationInflux	85	90	103	91	93	88	90	88	102	87	97	98	73	86	81	79
InfectionLevel	39	42	43	36	36	37	37	28	36	44	42	35	30	36	36	36
MONTH: 9	ALIVE: 1	INFECTED: 1	HEALTHY: 0	LAST MONTH DEATH: 3	LAST MONTH BORN: 1											
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
PopulationInflux	49	49	46	47	49	41	47	37	46	51	46	44	36	44	41	41
InfectionLevel	18	16	13	16	14	17	13	10	20	15	16	15	11	15	10	17
MONTH: 10	ALIVE: 0	INFECTED: 0	HEALTHY: 0	LAST MONTH DEATH: 1	LAST MONTH BORN: 0											
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
PopulationInflux	21	18	14	17	18	17	17	13	26	19	15	17	12	15	13	20
InfectionLevel	7	9	3	9	9	3	10	6	8	5	7	6	5	8	7	4
ALL SQUIRRELS DIED. SIMULATION ENDS.																

Figure 5: The output of a simulation. The parameters are: 34 healthy squirrels, 4 infected squirrels, 16 cell. The birth probability is 300.0 and the catch disease probability is 500.0.

3.1 Method 1

The first method is to look at the ALIVE, LAST MONTH DEATH, and LAST MONTH BORN presented in the output(Figure 5). If all squirrels have died before the terminal, these values should conform to the following formula(1):

$$ALIVE_0 + \sum LAST_MONTH_BORN = \sum LAST_MONTH_DEATH \quad (1)$$

The ALIVE_0 means the first month's ALIVE value-the initial squirrel number(including healthy and infected squirrels). The output of this simulation can pass this test.

SIMULATION STARTS.

MONTH: 1	ALIVE: 38	INFECTED: 4	HEALTHY: 34	LAST MONTH DEATH: 0	LAST MONTH BORN: 0														
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15			
PopulationInflux	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0			
InfectionLevel	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0			
cell 0 month 1 month_populationInflux = 134, month_infectionlevel = 16																			
MONTH: 2	ALIVE: 37	INFECTED: 3	HEALTHY: 34	LAST MONTH DEATH: 1	LAST MONTH BORN: 0														
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15			
PopulationInflux	134	122	134	125	127	107	115	123	147	127	121	133	142	122	136	113			
InfectionLevel	16	13	16	20	15	13	8	8	13	8	13	15	13	9	8	5			
cell 0 month 2 month_populationInflux = 111, month_infectionlevel = 9																			
MONTH: 3	ALIVE: 37	INFECTED: 10	HEALTHY: 27	LAST MONTH DEATH: 3	LAST MONTH BORN: 3														
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15			
PopulationInflux	245	217	257	238	243	220	238	240	279	245	234	251	288	220	242	226			
InfectionLevel	25	21	29	26	25	26	20	20	22	17	24	30	29	21	18	12			
cell 0 month 3 month_populationInflux = 124, month_infectionlevel = 54																			
MONTH: 4	ALIVE: 36	INFECTED: 15	HEALTHY: 21	LAST MONTH DEATH: 8	LAST MONTH BORN: 7														
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15			
PopulationInflux	369	326	368	360	376	335	353	379	413	366	372	371	417	350	354	335			
InfectionLevel	63	53	56	49	58	52	59	66	66	57	60	57	66	60	52	56			
cell 0 month 4 month_populationInflux = 115, month_infectionlevel = 70																			
MONTH: 5	ALIVE: 25	INFECTED: 18	HEALTHY: 7	LAST MONTH DEATH: 14	LAST MONTH BORN: 3														
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15			
PopulationInflux	350	299	334	348	357	334	341	348	366	332	365	324	358	308	328	317			
InfectionLevel	124	99	105	112	106	112	105	114	120	100	117	95	94	98	103	108			
cell 0 month 5 month_populationInflux = 43, month_infectionlevel = 36																			
MONTH: 6	ALIVE: 6	INFECTED: 5	HEALTHY: 1	LAST MONTH DEATH: 19	LAST MONTH BORN: 0														
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15			
PopulationInflux	282	241	265	277	282	276	272	274	285	272	299	246	260	251	287	250			
InfectionLevel	106	84	106	103	84	116	104	95	105	97	102	87	81	82	117	96			
cell 0 month 6 month_populationInflux = 10, month_infectionlevel = 9																			
MONTH: 7	ALIVE: 1	INFECTED: 1	HEALTHY: 0	LAST MONTH DEATH: 5	LAST MONTH BORN: 0														
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15			
PopulationInflux	168	142	158	170	159	172	162	148	157	163	171	140	135	133	187	155			
InfectionLevel	45	41	49	52	36	53	51	48	50	58	45	49	42	45	67	52			
cell 0 month 7 month_populationInflux = 1, month_infectionlevel = 1																			
MONTH: 8	ALIVE: 0	INFECTED: 0	HEALTHY: 0	LAST MONTH DEATH: 1	LAST MONTH BORN: 0														
CellId	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15			
PopulationInflux	54	48	61	57	51	68	59	57	57	70	57	54	52	55	79	61			
InfectionLevel	10	12	8	19	10	13	6	15	8	13	11	15	5	15	14	16			

ALL SQUIRRELS DIED.
SIMULATION ENDS.

Figure 6: The output of a simulation with printed cell information. The parameters are the same as that of Figure 5.

3.2 Method 2

The second method is to print a cell's single month's populationInflux and infectionLevel. In this way, we can test that the monthly updated information is correct. Figure 6 is the output with CELL_DEBUG = 1 and the tested cell Id is 0. The monthly update method is to sum up the past-3-month population and the past-2-month infection. From the Figure 6 it is easy to test the calculation of populationInflux and infectionLevel for cell 0. For example, the cell 0 populationInflux in month 4 can be calculated by summing up the cell 0 month_populationInflux of month 1, 2 and 3. In this way, we can prove that the values of monthly updated populationInflux and infectionLevel are convincing in this simulation.

4 Conclusions

This program can be successfully executed with different numbers of initial healthy and infected squirrels. It will appear an error when there are not enough processes provided. When the number of the squirrel is larger than 200, the program will appear the error message, and when all squirrels die, the program can be terminated immediately. Besides, the Macro definition 'CN' in file *simulation.h* can help in changing the number of cells.

However, it also has some issues needed to improve. One is that there are no user-define parameters provided and you can only modify parameters in *simulation.h*. Another issue is that there is not a framework constructed for this simulation. Therefore, it is not very portable to apply to other cases.